

## Project 1 - Operating Systems - New `getcnt()` syscall

DCC605 - Sistemas Operacionais - Universidade Federal de Minas Gerais

Group: Juan Marcos Braga Faria, Lucas Almeida Santos de Souza

### Tasks

1. Modify xv6 so the kernel keeps track of how many times each syscall has been called.
  - The `getcnt` syscall receives one integer parameter.
  - It returns the number of times the syscall with the given number has been called.
  - The user space declaration of `getcnt` should be `int getcnt(int)`
2. Implement a user-space program to call the syscall. The program should be named “`getcnt`” and receive a single integer as parameter corresponding to the target syscall.

Modified files:

- `kernel/syscall.c`
- `kernel/syscall.h`
- `kernel/sysfile.c`
- `kernel/proc.c`
- `kernel/proc.h`
- `user/user.h`
- `user/usys.pl`
- `user/getcnt.c`
- `Makefile`

### Implementation Choice

We may use different existing kernel functions as reference for `getcnt`. However that would depend on what processes the new functionality is meant to keep track of: the amount of times any system call has been called by a certain process, or by the whole system. On the first case, the `getpid` would be a good candidate as it looks up a variable local to the process. For the second case, maybe `uptime` would be better, as it has to acquire an external value. The function that has the job of finding and invoking the call, `syscall`, could be the primary candidate to do the incrementation task, albeit with a cost.

As the integer variable have limits, and the students have no data on how many times system calls are made during an expected system uptime - not to mention the performance overhead or the much longer uptime of servers-, the implementation chosen for `getcnt` will track how many times the system calls were called by each process.

It makes sense also as registering global values usually have limits imposed, such as the default number of registers for executed commands collected by history is 1000.

## xv6 modifications

### Data structure in proc.h

```
kernel\proc.h
@@ -84,24 +84,26 @@
84 84 // Per-process state
85 85 struct proc {
86 86     struct spinlock lock;
87 87
88 88     // p->lock must be held when using these:
89 89     enum procstate state; // Process state
90 90     void *chan; // If non-zero, sleeping on chan
91 91     int killed; // If non-zero, have been killed
92 92     int xstate; // Exit status to be returned to parent's wait
93 93     int pid; // Process ID
94 94
95 95     // wait_lock must be held when using this:
96 96     struct proc *parent; // Parent process
97 97
98 98     // these are private to the process, so p->lock need not be held.
99 99     uint64 kstack; // Virtual address of kernel stack
100 100     uint64 sz; // Size of process memory (bytes)
101 101     pagetable_t pagetable; // User page table
102 102     struct trapframe *trapframe; // data page for trampoline.S
103 103     struct context context; // swtch() here to run process
104 104     struct file *ofile[NOFILE]; // Open files
105 105     struct inode *cwd; // Current directory
106 106     char name[16]; // Process name (debugging)
107 107 +
108 108 + uint64 syscall_count[22]; // Array que contabiliza o número de vezes que cada syscall foi chamada
107 109 };
```

The chosen data structure was a simple integer array, `uint syscall_count`, to prioritise performance and access. It increases the length of a data structure that is created in every single process: `struct proc`.

**Why here?** As the chosen implementation keeps track of calls from each process, the data structure that contains each process's "metadata". It is initialized at the beginning of a process, and dies with it.

## Keeping track of calls in syscall.c

```
kernel\syscall.c
126 127 [SYS_link] sys_link,
127 128 [SYS_mkdir] sys_mkdir,
128 129 [SYS_close] sys_close,
130 130 + [SYS_getcnt] sys_getcnt,
129 131 };
130 132
131 133 void
132 134 syscall(void)
133 135 {
134 136     int num;
135 137     struct proc *p = myproc();
136 138
137 139     num = p->trapframe->a7;
138 140     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
139 141         // Use num to lookup the system call function for num, call it,
140 142         // and store its return value in p->trapframe->a0
143 143 +
144 144 +         // Incrementa o número de vezes que a syscall foi chamada
145 145 +         p->syscall_count[num-1]++;
146 146 +
141 147         p->trapframe->a0 = syscalls[num]();
142 148     } else {
143 149         printf("%d %s: unknown sys call %d\n",
144 150             p->pid, p->name, num);
145 151         p->trapframe->a0 = -1;
146 152     }
147 153 }
148 154
```

The data structure needs to be updated every time a system call is invoked. As `syscall_count` items matches the number of system calls, a simple jump using the call number as index should suffice: `p->syscall_count[num-1]++`.

**Why here?** The function responsible for validating and calling the call seems to be the best place to insert code to do these updates.

## Initializing data structure in `proc.c`

```
kernel\proc.c

105 105 // Look in the process table for an UNUSED proc.
106 106 // If found, initialize state required to run in the kernel,
107 107 // and return with p->lock held.
108 108 // If there are no free procs, or a memory allocation fails, return 0.
109 109 static struct proc*
110 110 allocproc(void)
111 111 {
112 112     struct proc *p;
113 113
114 114     for(p = proc; p < &proc[NPROC]; p++) {
115 115         acquire(&p->lock);
116 116         if(p->state == UNUSED) {
117 117             goto found;
118 118         } else {
119 119             release(&p->lock);
120 120         }
121 121     }
122 122 +
123 123 + // Inicializa syscall_count em zero para todas as syscalls.
124 124 + for (int i = 0; i < NELEM(p->syscall_count); i++)
125 125 +     p->syscall_count[i] = 0;
126 126 +
122 127     return 0;
123 128
124 129 found:
125 130     p->pid = allocpid();
126 131     p->state = USED;
```

According to the xv6 manual, `allocproc` is the function responsible for allocating a process's resources when it is initialized. It is the clear candidate for a code that is responsible for initializing data for `syscall_count`.

## Actual system function code in sysfile.c

```
kernel\sysfile.c
↑.... @@ -503,3 +503,18 @@ sys_pipe(void)
503 503     }
504 504     return 0;
505 505 }
+
506 +
507 + uint64
508 + sys_getcnt(void)
509 + {
510 +     int syscallID;
511 +     argint(0, &syscallID);
512 +
513 +     struct proc *p = myproc();
514 +
515 +     if(syscallID < 1 || syscallID > NELEM(p->syscall_count))
516 +         return -1;
517 +
518 +     int cnt = p->syscall_count[syscallID-1];
519 +     return cnt;
520 + }
```

`argint()` places the provided call number to the local `syscall` variable. It then acquires a pointer to the local `struct proc` through the `myproc()` interface, checking everything later. Grabbing the correct value is then a simple matter of reaching for the count value using the call number as index.

**Why here?** Naturally this file is where other similar system calls are located.

## Testing

```
user\getcnt.c
@@ -0,0 +1,22 @@
1 + #include "kernel/types.h"
2 + #include "kernel/stat.h"
3 + #include "user/user.h"
4 +
5 + int main(int argc, char **argv) {
6 +     if(argc != 2){
7 +         fprintf(2, "Usage: getcnt <syscall id>\n");
8 +         exit(1);
9 +     }
10 +
11 +     int syscall_num = atoi(argv[1]);
12 +
13 +     int cnt;
14 +
15 +     if((cnt = getcnt(syscall_num)) < 0){
16 +         fprintf(2, "getcnt: failed to get count\n");
17 +         exit(1);
18 +     }
19 +
20 +     fprintf(1, "syscall %d has been called %d times\n", syscall_num, cnt);
21 +     exit(0);
22 + }
```

The students used a series of simple tests. The user interface from Task 2, implemented in `user/getcnt.c` was used to check for the expected behaviors.

The students tried to verify the call doing the following:

- Call `getcnt 22` (the code of the `getcnt` system call itself) several consecutive times to check if the counter incremented with each call.
- Call `getcnt 2` on each execution to observe the behavior of the `exit` syscall (this syscall is called at the end of each command).
- Create a directory with `mkdir` and then call `getcnt(20)` (20 is the system call code for `mkdir`).

## Sources

<https://github.com/palladian1/xv6-annotated?tab=readme-ov-file>

<https://github.com/remzi-arpacidusseau/ostep-projects/tree/master/initial-xv6>

<https://moss.cs.iit.edu/cs450/mp1-xv6.html>

<https://pdos.csail.mit.edu/6.828/2023/xv6/book-riscv-rev3.pdf>